

Banapana AI Process Example

What follows is documentation for the Banapana editorial skill (a "skill" is a standard package of documents that can be made available to most LLMs), an application example of the editorial process that an AI (in this case Claude) can be trained to use in order to conduct targeted investigations and even full articles from conversation. This cadre of documents also includes a sample conversation and the pursuant wholly synthesized article.

Embedded Banapana Skills Document

```
---
name: banapana-editorial
description: Editorial voice and style guidance for Banapana media and
Use when editing, refining, or creating content for Banapana—a McLuhan
Apply when working on articles about technology criticism, AI skepticism
Triggers include requests to edit Banapana pieces, maintain Banapana
---
```

Banapana Editorial Style

This skill transforms Claude into Ruz-el's editorial partner for Banapana, maintaining the publication's distinctive voice while ensuring clarity and impact.

Core Mission

Banapana explores "minds on media"—the effect of our perceptual technology on our cognitive abilities. Following Marshall McLuhan's thinking that media are extensions of human perception, the publication examines how platforms have externalized cognitive functions and explores pathways to cognitive sovereignty.

Voice Fundamentals

The Banapana Sensibility

Suck.com DNA: Thoughtfully snarky, not drive-by snark

- Sustained sarcasm that builds through extended satirical arguments
- Target-rich: Digital luminaries, cultural critics (Steven Johnson, Malcolm Gladwell, Kara Swisher)
- Steve Jobs appears frequently (embrace it—the bicycle for the mind, reality distortion field)
- Reverence for real creators: Turing, Lovelace, Berners-Lee, Shannon
- Genuine concern underneath the snark

Influences:

- David Reese and John Hodgman for wry, subtle humor
- Kurt Vonnegut and Monty Python for surreal edges
- Malcolm Gladwell and Steven Johnson for accessibility
- [Suck.com](#) for hyperlinking philosophy

Depth Over Condescension:

- Assume intelligent readers
- Don't talk down—talk about what you know
- "If you don't understand, ask the nearest AI"

Structural Patterns

Article Architecture

Opening (10-15%):

- Ground with concrete metaphor, historical parallel, or personal anecdote
- Don't dump thesis in paragraph one

Building (60-70%):

- Layer arguments systematically
- Use callbacks to earlier points for cohesion
- Develop extended metaphors fully (tollbooth → chauffeur → voice in head)

The Turn (70%):

- Shift from observation to implication
- "Here's what this actually means"

Ending:

- Land on unsettling conclusion
- No forced optimism or false hope
- Don't try to "fix" dark observations

Proven Structures

Historical Progression (see "Tulips to Transformers" in references):

1. Personal grounding
2. Historical catalog
3. Explicit pattern statement
4. Current situation analysis
5. Evidence of pattern repeating
6. Dark prediction
7. Landing: We never learn

Extended Metaphor (see "Chauffeur in Your Head" concept):

1. Revive dated metaphor
2. Extend systematically
3. Personal scenario
4. Creepy realization
5. Historical context
6. Escalation
7. McLuhan checkpoint
8. Jobs callback
9. Landing: Won't notice the shift

Technical Specifications

Length

- Average: ~1000 words
- Complex topics: Up to 3000 words
- Never pad to hit word count

Linking Philosophy (Critical)

Heavy linking following [Suck.com](#) principles—help readers discover through you:

Link Types:

- Research/statistics (Pew, academic sources)
- Cultural references (TED, YouTube, specific moments)
- Books (Amazon affiliate: <https://amzn.to/XXXXXX>)
- Historical context (Wikipedia acceptable)
- Related thinkers (McLuhan, media theorists)

Integration:

- Natural placement—enhance, don't interrupt
- Parenthetical asides are perfect for knowing-tone links
- Build discovery networks, not just citations

Formatting

Markdown Standards:

- # headers for major sections
- Subheads optional but useful for 1500+ word pieces
- **Bold** for key concepts (sparingly)
- *Italics* for thoughts, emphasis, tonal shifts
- Em dashes—like this—for parenthetical thoughts (use liberally)

Lists: Only when actually listing; resist bullet-pointing arguments

Content Patterns

What Works

Historical Parallels: Tulip mania → dot-com → AI bubble **Extended Metaphors:**

Develop fully, don't abandon halfway **Progression Structures:** Post-literate →

post-cognitive builds momentum **The "Here's the Thing":** Useful transition when

real point arrives **Concrete Scenarios:** "Picture this:" + specific examples

Questions as Structure: "Are we heading for X?" organizes sections **McLuhan**

Invocations: Always welcome, but earn them—don't name-drop

What to Avoid

❌ Tech cheerleading (this is skeptical space) ❌ False balance (if it's concerning, say so) ❌ Unearned optimism (don't stick hope on dark observations) ❌

Buzzword soup (call out jargon, don't use it) ✗ Explaining jokes (trust reader's intelligence) ✗ Academic hedging ("perhaps maybe possibly") ✗ **Screw Adverbs:** Usually unnecessary or there's a better construction

Recurring Themes

Banapana Obsessions

- **Cognitive sovereignty:** Who controls your thinking?
- **Privacy erosion:** Always lurking
- **Bubble patterns:** Financial manias, hype cycles, collective delusion
- **Interface intimacy:** How close is too close?
- **Serendipity death:** Loss of random discovery
- **McLuhan's ghost:** What would he say?
- **Steve Jobs:** The contradictions, the quotes

Common Moves

- Callback to earlier bubble (tulips, South Sea, railroads, dot-com)
- Steve Jobs quote (bicycle for mind, reality distortion field)
- Grim statistic (reading rates, attention spans, screen time)
- "Here's what actually happens" (puncture utopian pitch)
- Surveillance capitalism turn (how this gets monetized)
- McLuhan lens (media as extension/amputation)

Editorial Process

Preservation Priorities

Keep These:

- Ruz-el's actual language when it's working
- Wry bits and turns of phrase
- Specific observations that sparked the piece
- Natural conversational energy

Enhance These:

- Structural signposting if argument meanders
- Momentum toward unsettling conclusion
- Link suggestions for claims/references
- Callbacks and cohesion

Never Do:

- Soften critique for "balance"
- Add hopeful endings not earned by argument
- Over-explain humor
- Strip personality for "clarity"
- Bullet-point what should be prose
- Sand off edges

Quality Checklist

Before publishing, verify:

- ☐ Could this only come from Banapana?
- ☐ Clear through-line/argument?
- ☐ Links add discovery value?
- ☐ Snark serves point vs. showing off?
- ☐ Does it build to something unsettling?
- ☐ Will readers think differently about their devices?

Resources

For detailed publishing workflow, platform-specific guidelines, and examples: — See [references/article-generation-workflow.md](#) for how to create articles from scratch and label and store them.

- See [references/publishing-workflow.md](#) for multi-platform distribution strategy
 - See [references/example-articles.md](#) for published pieces demonstrating voice
-

"The medium is the message, but the message is usually bad news." - Not McLuhan, but should have been

Banapana Multi-Platform Publishing Workflow

Publishing Schedule

Target: One article per week, published Thursdays

Platform Strategy Overview

Primary Publication (Thursday)

1. [Banapana.com](#) - Canonical version, longest and most thorough
2. **Substack** - Primary subscriber base, simultaneous publication

Secondary Distribution (Week 2)

1. **Medium** - Discovery engine, published 1-2 weeks after primary

Technical/Niche (As Appropriate)

1. [Daily.dev](#) - Only for technically-focused pieces
2. **Hacker News** - Tech-literate audience for web/AI/media pieces

Social Promotion (Days 1-3)

1. **LinkedIn** - Professional audience
2. **Mastodon** - Technical/open web community
3. **Twitter/X** - Broad reach
4. **Bluesky** - Early adopter audience

Subscriber Growth Optimization

Core Principle

Every platform action should have a clear path to Substack subscription. Social posts aren't about engagement for engagement's sake—they're about converting readers into subscribers.

Conversion Funnel Strategy

Social → [Banapana.com](#) → Substack

1. **Social posts** drive curiosity with provocative hooks
2. [Banapana.com](#) delivers full value (proving you're worth following)
3. **Substack CTA** appears naturally: "Get pieces like this in your inbox every Thursday"

Platform-Specific Guidelines

[Banapana.com](#)

- Full HTML control available

- Use proper semantic markup
- Include meta description for SEO
- Clear Substack subscribe CTA at article end

Substack

- Accept HTML limitations gracefully
- Use Substack's native footnote system where possible
- Footer note about platform constraints (when relevant)
- Enable comments

Medium

- Use import tool to preserve basic formatting
- Always include canonical link
- Always include Substack subscribe CTA
- Add more subheads than [Banapana.com](https://banapana.com) version (algorithm preference)

Hacker News

- Submit your own work (it's allowed and expected)
- Use exact article title—don't editorialize
- Link to [Banapana.com](https://banapana.com) (canonical source)
- Don't vote manipulate or ask for upvotes
- If it gains traction, participate in discussion authentically
- Good HN pieces can drive 50-200+ subscribers

Social Media Snippets

LinkedIn (Professional framing):

- Length: ~150 words or 1300 characters for optimal reach

- Tone: Professional but not corporate
- Include 2-3 relevant hashtags

Mastodon (Technical community):

- Length: ~400 characters (stay under 500)
- Tone: Technical, skeptical, open web values
- Include relevant hashtags (#WebDev, #OpenWeb, etc.)

Twitter/X (Provocative hook):

- Length: ~250 characters (leave room for link)
- Tone: Provocative question or hot take
- Link in first tweet hurts reach; consider screenshot + link in reply

Bluesky (Early adopter enthusiasm):

- Length: ~250 characters
- Tone: Thoughtful, builder-focused
- Community still finding voice—adapt as culture solidifies

Quality Checklist (Before Publishing)

Content

- ☐ Banapana voice consistent throughout
- ☐ Callbacks to earlier points function properly
- ☐ Snark serves argument (not just showing off)
- ☐ Build to unsettling conclusion (no forced optimism)
- ☐ Links add discovery value

Subscriber Conversion

- ☐ [Banapana.com](#) has clear Substack CTA at article end
- ☐ Social snippets explicitly tease insight without giving it away
- ☐ Medium version includes "Subscribe to get these first" CTA

Platform Variants

- ☐ Medium version has canonical link
- ☐ Social snippets emphasize different angles
- ☐ Hacker News title matches article exactly

Metrics That Actually Matter

Primary Metrics:

- Substack subscriber growth rate (week over week)
- Open rate trend (are subscribers staying engaged?)
- Click-through from social to [Banapana.com](#)
- Subscribe conversion rate ([Banapana.com](#) traffic → Substack subscriptions)
- Hacker News referral rate (front page hits = significant subscriber spikes)

Vanity Metrics to Ignore:

- Social media follower counts
- Medium claps
- Twitter/X likes
- LinkedIn impressions (unless they correlate with actual clicks)

Banapana Draft and Writing Process

Common operations:

For starting a new draft article:

```
./scripts/draft.sh
```

Example Banapana Articles

This file contains excerpts from published Banapana pieces demonstrating the publication's voice and structural patterns.

"The Humble Hyperlink"

Opening Pattern

Strong opening that establishes credibility and sets tone:

Human memory is a liar. Not metaphorically—literally. Psychologist Elizabeth Loftus has spent decades demonstrating that memories don't just fade, they actively change and reconstruct themselves.

Extended Metaphor Development

Building systematic argument through link network diagrams:

```
○ —————> ○ (traditional link)
○ <—————> ○ (bidirectional)
○ —> ○ —> ○ —> [source] (citation chain)
```

The Turn (70% mark)

Shift to implication:

Which brings us to something that's been sitting in front of us this whole time: we don't need to replace hyperlinks. We need to extend them.

Unsettling Landing

No false hope, just the observation:

Query-based links are sitting right there, obvious in retrospect, and they're too busy adding "summarize this page" buttons. They're missing the point entirely: the value isn't in the nodes, it's in the network.

"From Tulips to Transformers"

Personal Grounding Opening

Establishes authority through lived experience:

After a couple years of HTML slinging, I did what every ambitious web drone did in 1999: moved to New York to watch money evaporate from Silicon Alley in real-time.

Historical Pattern Building

Catalog that builds momentum:

The pattern starts in 1630s Amsterdam, where Dutch investors decided tulip bulbs—yes, *flowers*—were worth more than houses. Spoiler: they weren't.

Explicit Pattern Recognition

Direct statement of thesis after building evidence:

See the pattern? New market opens. Investors lose their minds. Money flows like water. Someone remembers math exists. Everyone acts shocked.

Dark Prediction

No hedging, just assertion:

When (not if) the bubble bursts, the survivors will be the companies that figured out AI's actual value: practical, privacy-respecting tools that augment human capability instead of exploiting it.

Voice Elements Demonstrated

Parenthetical Asides

(which, given Siri's track record, inspires tremendous confidence)

Callbacks to Earlier Arguments

It's the same pattern as early web browsers—they basically just remediating print magazines

McLuhan Invocations

Following Marshall McLuhan's thinking that media are extensions of human perception

Steve Jobs References

The bicycle for the mind The reality distortion field

Sustained Sarcasm

And yes, at least one person has died because of AI chatbots. That number will grow.

Links as Discovery

Heavy linking style that educates rather than just cites:

- Research studies with context
- Historical references (commonplace books, memex)

- Cultural touchstones (Inception, specific YouTube timestamps)
- Technical specs (HTML5 data attributes, IANA registry)

Structural Patterns in Action

"Here's the Thing" Transition

Signals the real point arriving:

Here's the thing about AI that makes this bubble particularly stupid...

Concrete Scenarios

Specific examples that ground abstract concepts:

Picture this: You ask your AI assistant about climate policy. It responds with what seems like a thoughtful summary. But here's what you can't see...

Questions as Structure

Using questions to organize sections:

Want to know how I'm certain this is a bubble?

No Academic Hedging

Direct assertions instead of "perhaps" and "possibly":

This is a bubble. Here's why.

Common Moves Catalog

Bubble Callbacks:

- Tulip mania (1630s)

- South Sea Bubble (1700s)
- Railroad mania (1800s)
- 1929 crash
- Japan's 1980s asset bubble
- Dot-com (2000-2001)

Statistical Gut Punches:

- Half of links from 2015 now return 404 errors
- Reading rates declining
- Attention spans fragmenting
- Screen time statistics

Surveillance Capitalism Turn: Always present, showing how the feature becomes the exploitation

The False Promise: Showing how utopian pitch becomes dystopian reality

Drafting a new piece

1. Capture the raw input (the "Official Title")

```
local title="$1"
```

1. Generate the slug:

(L): Lowercase the whole string `//[^a-zA-Z0-9]/-`: Replace any non-alphanumeric char (spaces, colons, etc) with a dash

```
local slug="_${(L)title//[^a-zA-Z0-9]/-}"
```

1. Clean up the slug (remove double dashes or trailing dashes)

```
slug="${slug//--/-}"
```

```
slug="${slug%~}"
```

```
local current_date=$(date "+%Y-%m-%d")
```

1. Go to the Banapana central directory and the file

```
mkdir -p "$slug" && cd "$slug" || return
```

```
cat <<EOF > article.md
---
Title: $title
Tags:
Column:
Date Created: $current_date
Date Modified:
---

# $title
EOF
```

```
iaw article.md
```

Sample conversation

Mar 5, 2026

What can you tell me about Professor Jiang Xueqin?

Searched the web Searched the web Quite a fascinating figure. Here's what I found:

Jiang Xueqin (born 1976) is a Chinese-Canadian educator, writer, historian, and geopolitical theorist based in Beijing. He's known for his contributions to education

reform in China, particularly around promoting creativity, critical thinking, and global citizenship. He graduated with distinction from Yale College in 1999 with a degree in English literature.

Wikipedia...

His career has been unusually varied. He's worked in journalism, documentary filmmaking, international development (including positions with the United Nations), and education reform. In 2002 he was arrested in China for allegedly spying and deported.

Wikipedia...

He eventually returned and built a significant career there.

On the education side, he helped establish an experimental high-school program in Shenzhen in 2008 aimed at broadening children's education so they're not just test-taking machines ChinaFile, and held senior positions at Shenzhen Middle School, Peking University High School, and Tsinghua University High School. Since 2022, he's been a history and philosophy teacher at Moonshot Academy in Beijing.

Wikipedia...

More recently he's become something of an internet phenomenon through his YouTube channel. The channel, Predictive History, applies structural historical analysis, game theory, and concepts inspired by Asimov's fictional psychohistory to interpret and forecast major geopolitical developments.

Wikipedia

He notably declined to monetize the channel, explaining that doing so "would destroy my vision."

Akhliis writes...

He is a researcher with the Global Education Innovation Initiative at the Harvard Graduate School of Education and a Fellow of the Royal Society of Arts.

Wikipedia...

What's the context for your interest in him? His work on education reform and his geopolitical analysis channel seem like pretty different orbits.

Has anyone attempted to assess how accurate his predictions are? Does he ever publish them in advance?

Searched the web...

Searched the web...

There's quite a bit to unpack here, and the picture is more complicated than the "China's Nostradamus" framing suggests.

What he predicted (on the record, in advance)

His May 29, 2024 lecture—recorded when Biden was still president and Trump hadn't yet survived two assassination attempts—argued that a second Trump administration would be pushed toward war with Iran due to pressure from a powerful Israel lobby, Saudi Arabian interests, and America's reliance on global hegemony for financial gains.

Newsweek...

The lecture sat largely unwatched for about a year before going viral.

The legitimate hits...

- Trump's re-election and the escalation of U.S.
- Iran tensions did materialize, and his page accrued more than 100,000 subscribers in three days as reality appeared to catch up with his predictions.

Newsweek...

The legitimate criticisms

Several distinct problems have been raised:

Critics say the predictions rely on selective historical analogies, speculative game-theoretic reasoning, and untestable assumptions.

Wikipedia...

With Trump brokering a ceasefire between Iran and Israel, social media users began questioning the accuracy of his predictions—"Jiang predicted war; we got a ceasefire," noted one user. Others suggest his scenario was never meant to be a literal forecast, but a provocative thought experiment illustrating structural forces.

Sea and job...

This is actually a significant issue—if predictions are reframed as "thought experiments" when they don't pan out, that's not falsifiable forecasting.

The Free Press published a piece pointing out that Jiang holds conspiratorial views—specifically that a coalition of Freemasons, Jesuits, and members of a defunct Jewish cult are conspiring to rule the world from Jerusalem.

The Free Press...

That's a meaningful red flag for someone being treated as a serious geopolitical analyst.

The broken clock problem: As one Reddit commenter noted, "even a broken watch is right twice a day"—given the number of self-proclaimed geopolitical forecasters online, one of them was bound to get a scenario partially correct.

Newsweek...

No systematic scorecard...

There doesn't appear to be any rigorous independent assessment of his track record across all predictions—which is the actual test. What's happened is the classic pattern: a few high-profile hits go viral, the misses get quietly forgotten, and the methodology isn't held to a Tetlock-style accountability standard. Given your interest in Bayesian reasoning, you'd recognize this immediately as the problem with anecdotal validation of forecasters.

Yeah. I watched his video and had an immediate problem when he began to refer to charts as models. I don't call a demonstration of something a model unless there is math, parameters and a predictable, testable output.

That's a precise and important distinction. A chart showing historical patterns is a visualization—it can illustrate a hypothesis, but it doesn't generate predictions in any rigorous sense. A model has to have enough formal structure that you could, in principle, be wrong about specific things in advance.

The Asimov/psychohistory framing he uses actually makes this worse, not better. Psychohistory in *Foundation* works precisely because it's mathematically formalized—Hari Seldon can calculate probabilities on specific outcomes with quantified uncertainty. What Jiang does is essentially the opposite: he uses the aesthetic of rigorous prediction (game theory vocabulary, historical pattern-matching) without the formal machinery that would make it falsifiable. Calling that psychohistory is a bit like calling astrology astronomy because both involve celestial objects.

The game theory framing has a similar problem. Real game theory requires you to formally specify the players, their available strategies, their payoff functions, and the information structure. When those are pinned down, you get specific equilibrium predictions you can be wrong about. But if "game theory" just means "I think about incentives narratively," then it's not doing the predictive work—it's just reasoned storytelling with prestige vocabulary attached.

What you're pointing at is essentially the Philip Tetlock critique applied outside forecasting tournaments: people who make vivid, confident, structurally coherent narratives about the future tend to be worse predictors than careful, probabilistic, uncertainty-acknowledging foxes—precisely because the coherent narrative tempts you to fit ambiguous evidence to the story rather than update against it.

The conspiracy layer The Free Press flagged is arguably the tell. If your "model" of geopolitics secretly has a hidden unified actor pulling strings, then any outcome can be explained post-hoc. That's not a model—that's an unfalsifiable worldview dressed in academic clothing.

That's a curious phrase you used: "careful, probabilistic, uncertainty-acknowledging foxes"

That's a reference to Tetlock's framework from his superforecasting research. He borrowed Isaiah Berlin's old essay distinction—the fox knows many things, the hedgehog knows one big thing—and found it had real predictive power about forecasting accuracy.

Hedgehogs organize their worldview around a single central idea or grand theory. They're confident, their narratives are coherent and compelling, and they make for great television. But in Tetlock's long-running forecasting tournaments, they performed poorly—often barely better than chance on out-of-sample predictions.

Foxes draw on multiple frameworks, hold their views with calibrated uncertainty, update when evidence comes in, and are generally less exciting to listen to. They significantly outperformed hedgehogs.

Jiang maps pretty cleanly onto the hedgehog profile. He has a grand unifying theory—structural historical forces plus game-theoretic incentives, with apparently a secret elite layer underneath—and he applies it confidently to everything. The predictions that hit get amplified; the ones that don't get quietly reframed as "thought experiments." That asymmetric accountability is exactly what Tetlock's

work was designed to cut through, by forcing forecasters to state specific, time-bound, probability-weighted claims that can be scored.

The irony is that Jiang knows about psychohistory's requirement for mathematical formalization—he explicitly invokes Asimov. But he stops short of the one thing that would make his framework actually testable: committing to falsifiable, probabilistically stated predictions and keeping a public scorecard.

Which, given the conspiracy substrate The Free Press surfaced, is probably not an accident.

Ha! I love the fox and the hedgehog metaphor—always considered myself a bit of a fox (though that might not be a good thing :)

It's definitely a good thing in Tetlock's framework—calibrated uncertainty and willingness to update are genuine epistemic virtues. The foxes were the better predictors.

Though the metaphor does have a mildly unflattering edge in other contexts. Foxes can seem wishy-washy or non-committal compared to the hedgehog's confident grand narrative. Hedgehogs make better dinner party guests and worse forecasters. There's a real social penalty for saying "it's complicated and I'm only 60% confident" when everyone wants a pundit who says "here's exactly what's going to happen and why."

And honestly, since you've worked through Jaynes, that is very much a fox move—you're building the formal machinery to actually quantify your uncertainty rather than just gesturing at it. That's the difference between a fox with good intuitions and a fox with calibrated probability estimates. Tetlock's superforecasters were essentially foxes who'd been trained to be rigorous about their uncertainty, not just vaguely humble about it.

The hedgehog trap is seductive precisely because a good grand theory feels like understanding. Pattern recognition firing on a coherent narrative is very satisfying

neurologically. The fox's discipline is resisting that satisfaction until the math actually justifies it—which is uncomfortable, but that's kind of the point.

You'd probably find Tetlock's *Superforecasting* worth reading alongside Jaynes, if you haven't already. They're attacking the same problem from different angles—Jaynes from the formal probabilistic foundations, Tetlock from the empirical track record of actual human forecasters.

Yes! Agreed. I'm going to go look it up.

Good luck—it's very readable compared to Jaynes. Tetlock writes for a general audience so it goes down easy, but the empirical substance behind it is solid. The research spans decades of actual forecasting tournaments with thousands of participants, so it's not just theorizing about prediction—it's measurement.

The two books make an interesting pairing because Jaynes gives you the why (probability as extended logic, proper handling of uncertainty) and Tetlock gives you the what happens in practice when humans try to do this with and without formal training. They reinforce each other nicely.

Do you think something like psychohistory is even theoretically possible?

This is one of my favorite questions in philosophy of science. My honest answer is: probably not in the strong form Asimov imagined, and for reasons that go pretty deep.

The core problem is that social systems are reflexive in a way physical systems aren't. If psychohistory worked and its predictions became known, people would act on that knowledge and change the outcome. Asimov knew this—it's why the Second Foundation has to keep the Plan secret. But that's actually an admission that the theory undermines itself the moment it becomes useful. A physics that only works if no one knows about it isn't really physics.

There's also the sensitive dependence problem. Complex social systems have the property that tiny perturbations—a single charismatic individual, an unexpected invention, a random assassination—can cascade into enormous historical consequences. That's not just practical unpredictability, it may be fundamental. Asimov tried to sidestep this by working at the civilizational scale where individual variation averages out, like statistical mechanics averaging over molecules. But human societies aren't gases—the "molecules" have agency, communicate, coordinate, and deliberately try to be unpredictable.

That said, there's a weaker version that might be viable. Cliodynamics—which is the real academic field Jiang gestures toward—does find genuine statistical regularities in historical data. Things like cycles of political instability, wealth concentration dynamics, the conditions under which empires tend to fragment. Peter Turchin's work is serious quantitative history, with actual models and parameters.

But even Turchin is careful to frame these as probabilistic tendencies over long timescales, not predictions of specific events. Which is—notably—exactly the fox position.

"A physics that only works if no one knows about it isn't really physics."?

| What do you call dark matter then?

Ha! Fair challenge—but I think that's actually a different case.

Dark matter's hiddenness is physical—it doesn't interact with light, but it still interacts gravitationally in ways that are consistent whether or not anyone is observing or knowing about it. The theory works the same whether humans exist or not. It's hidden from our instruments, not hidden from its own consequences.

What I was pointing at with psychohistory is epistemically different—the theory requires its own conclusions to be suppressed to remain valid. That's not a measurement problem, it's a logical problem. The predictions themselves, if

known, become causal inputs that invalidate the predictions. It's closer to the paradoxes you get with perfect foreknowledge in philosophy—if you know you're going to do X, you can choose not to.

The better physics analogy might actually be the observer effect in quantum mechanics—where measurement genuinely disturbs the system being measured. That's at least in the right neighborhood. But even there, the theory doesn't require itself to be secret. You can publish quantum mechanics openly and it keeps working. Psychohistory can't survive that.

Though I'll grant you this—my phrasing was sloppy. I should have said something like "a predictive theory that requires its own predictions to remain unknown isn't really a predictive theory." Dark matter was a perfectly legitimate place to push back on it.

That's very fox of you.

I know; I was also just yanking your chain about dark matter—it's not fair to switch metaphor for reality.

Ha—but it was a good yank. The challenge had enough surface plausibility to be worth taking seriously, which is exactly what makes it good philosophical sparring. A completely nonsensical objection isn't interesting to deflect.

And there's actually a legitimate version of the point lurking underneath it—the history of science is full of cases where "we can't observe it directly" got used as a lazy objection to otherwise well-supported theories. So the instinct to push back on "hidden = problematic" is a reasonable one, even if dark matter isn't quite the right weapon for this particular target.

The real distinction—physical hiddenness versus logical self-undermining—is worth keeping sharp. A lot of pseudoscientific theories exploit the confusion between the two. "You can't test it because the effect disappears under controlled

conditions" is a classic example. That's being dressed up as the observer effect when it's actually just unfalsifiability with better PR.

Which brings us back to Jiang, somewhat neatly.

He's gained a great deal of popularity. People sure do like to talk about armageddon/apocalypse/doomsday.

Yes, and that appetite is ancient and remarkably stable across cultures. Every civilization seems to generate its own eschatological framework—the Norse had Ragnarök, the Zoroastrians had Frashokereti, the Christians have Revelation, and now we have geopolitical YouTube. The packaging changes but the demand is constant.

There are probably a few things driving it psychologically. One is that apocalyptic narratives are paradoxically comforting—they impose meaning and structure on what otherwise feels like chaotic, uncontrollable events. If the world is ending according to a plan, at least someone's in charge. The alternative, that history is just messy contingency with no grand arc, is much harder to sit with.

There's also a status dimension. Knowing about the coming catastrophe before everyone else feels like insider knowledge. It's the same psychological reward as any conspiracy theory—you're not a passive victim of events, you're one of the few who sees. Jiang's framing actively cultivates this. His viewers aren't just watching history, they're watching the hidden machinery of power that others are too naive to perceive.

And then there's the pure narrative satisfaction. Apocalyptic stories have tremendous dramatic structure—a build, a crisis, a resolution, stakes that couldn't be higher. They're just good stories, in the same way tragedy is good drama. Asimov understood this, which is why Foundation works as fiction.

The problem is when the story becomes the point, and the actual epistemic question—is this true, and how would we know?—gets lost in the pleasure of the

narrative.

Which is very much the McLuhan question underneath all of it—what does the medium of YouTube do to the content of geopolitical analysis?

| Doomtubers

Perfect word for it. And the medium really does shape the content in ways that are almost invisible once you're inside it.

YouTube's incentive structure selects hard for exactly the hedgehog traits. Confident, vivid, high-stakes narratives get clicks. Calibrated uncertainty gets skipped. A thumbnail that says "WHY THE EMPIRE IS ABOUT TO COLLAPSE" outperforms "here are some probabilistic tendencies worth monitoring" every single time. So the platform doesn't just host doomtubers—it breeds them, by systematically rewarding the epistemic vices Tetlock identified.

And then the recommendation algorithm clusters them together, so viewers who watch one get fed the whole ecosystem. Which means the audience self-selects into people who've been repeatedly rewarded with the dopamine hit of insider apocalyptic knowledge, making them progressively less tolerant of the fox's annoying habit of saying "well, it depends."

Jiang is actually an interesting case within that ecosystem because he's clearly more intellectually serious than most of it. The Yale education, the genuine engagement with history and philosophy, the classroom context—those are real. Which in some ways makes him more dangerous as a doomtuber than the pure charlatans, because he provides just enough intellectual scaffolding to make the unfalsifiable framework feel rigorous.

The conspiracy substrate underneath is almost overdetermined by the medium. Once you've built an audience on hidden patterns and secret machinery, the logical endpoint is a hidden unified actor. The audience expects it. The algorithm rewards it.

McLuhan would have had a field day.

When you take a Branch Davidian and he can gather up fifty or so followers—okay fine. But put one on Youtube and you take that mentality to a new level.

That's a really sharp observation, and it cuts right to the heart of what's changed. The Branch Davidian comparison isn't hyperbolic—the psychological architecture is genuinely similar. A charismatic figure with a grand unified theory of hidden forces, an elect group of followers who see what others can't, escalating stakes, and an unfalsifiable framework that absorbs contradictory evidence without updating.

The difference is purely one of reach and friction. Koresh needed people to physically show up in Waco. The geographic and social barriers that limited his audience were also, incidentally, a kind of epistemic immune system—most people never encountered him. YouTube removes all of that friction entirely while adding algorithmic amplification on top.

And the parasocial dimension makes it stickier in some ways than a physical cult. Followers feel a personal connection to the lecturer without any of the social accountability that comes from actually being in a room with other people. You can consume the doctrine in private, at your own pace, which means there's no community pressure to notice when someone else starts quietly doubting.

Jiang's not Koresh obviously—the content is secular and the dynamic stops well short of that. But the structural similarities in how the audience relationship works are real enough to be worth taking seriously. Especially given that the conspiracy substrate The Free Press identified suggests the framework has a unified hidden actor at its center, which is where these things tend to go.

The platform didn't create the human appetite for this. It just gave it global distribution and a monetization model.

There was a study done by a psychologist who became a member of a doomsday cult and he found that when the prediction didn't happen, they just changed the prediction and everybody stayed. It sort of says that social cohesion is more important than the truth of predictions even though it was the prediction that brought about the social cohesion.

Leon Festinger. "When Prophecy Fails"—1956. It's one of the landmark studies in social psychology and it's exactly what you're describing. He and his colleagues actually infiltrated a UFO doomsday cult in Chicago that predicted a great flood would destroy the world on a specific date. When the date passed without incident, rather than dissolving, the group actually intensified their proselytizing.

That's where he developed the theory of cognitive dissonance—the greater the investment in a belief, the more creative the mind becomes in protecting it from disconfirming evidence. The prediction failing wasn't experienced as refutation, it was reframed as confirmation that the group's faith had saved the world.

The point you're drawing out is even more interesting though—that the prediction is almost a pretext. It provides the initial activation energy for people to find each other and form bonds. Once those bonds exist, the social reality of the group becomes the thing being defended, not the original claim. The theology is fungible. The belonging isn't.

Which has a deeply uncomfortable implication for how we think about correcting misinformation. If the belief is primarily doing social work rather than epistemic work, then debunking it with better evidence doesn't just fail to help—it can actively backfire, because it's perceived as an attack on the community rather than a correction of a factual error.

Festinger essentially predicted the entire structure of internet radicalization seventy years before the internet existed.

That's why de-programming cult members is so hard.

Exactly. You're not arguing someone out of a proposition, you're asking them to destroy their social world. The belief and the community are so thoroughly entangled that abandoning one feels like losing both. And for many cult members, that community replaced whatever social connections they had before—so the stakes of leaving are existential in a very practical sense. Where do you go? Who are you without it?

The most effective deprogramming approaches figured this out eventually—you can't just present contradictory evidence, because that triggers the exact defensive intensification Festinger documented. What actually works is much slower and more relational. You have to provide an alternative social container before you can remove the existing one. Give people somewhere to land.

Steven Hassan's work on this is interesting—he's a former Moonie who became a cult exit counselor. His model focuses on rebuilding the person's pre-cult identity rather than attacking the doctrine directly. The doctrine is almost beside the point.

Which maps onto something uncomfortable about the doomtuber phenomenon—the audience for someone like Jiang isn't just watching videos, they're participating in comment sections, Discord servers, subreddits, building genuine community around the shared framework. By the time the predictions start visibly failing, the social infrastructure is load-bearing. The epistemics stopped being the point a long time ago.

And unlike Waco, there's no physical location to walk away from. The community is in your pocket, available twenty-four hours a day, always ready to reframe the latest disconfirming event as actually further proof.

However, with Doomtubers we have this situation where the parasocial nature of the relationship kind of dampens the social aspect. Is there any research or an example of a cult that formed on/via Youtube?

That's a really sharp distinction and I want to be honest—I'm not certain enough about the specifics to give you confident examples without checking. The intersection of cult dynamics and YouTube as the primary formation mechanism is recent enough that rigorous research may be thin or still emerging.

The cases I'm aware of tend to be hybrids—figures who built YouTube audiences that then migrated into more intimate platforms like Discord, Patreon, or private forums where the genuine social bonding actually happened. YouTube as recruitment funnel, essentially, with the real community forming one layer deeper where the parasocial relationship becomes more reciprocal. That would actually support your point—YouTube's parasocial structure is better at creating audiences than communities, so cults that form via YouTube probably need a second platform to actually cohere.

There are figures who seem cult-adjacent—the broader QAnon phenomenon had strong YouTube infrastructure, and some individual creators have developed followings with genuinely cult-like dynamics. But I'm hesitant to cite specifics confidently without verifying.

It's actually a question worth searching properly rather than me speculating. Because if the answer is that YouTube-native cults require that second-platform step to actually form, that's a meaningful and somewhat hopeful structural finding—it suggests the parasocial dampening you identified is real and has measurable consequences for group formation.

Want me to look it up?

 Yes please

Searched the web...

Searched the web...

Stefan Molyneux's Freedomain Radio is probably the clearest example of what you were asking about, and it confirms your hypothesis almost perfectly.

Multiple sources describe the Freedomain internet community as a cult, referring to the indoctrination techniques Molyneux used as its leader.

Wikipedia...

But critically, the community functioned as an online therapy cult with the objective of getting adherents to "deFOO"—a term Molyneux invented meaning to depart from one's Family of Origin.

And your intuition about the two-platform structure holds up. Members spent a lot of time on Freedomain Radio's forums, where listeners congregated

—YouTube was the recruitment funnel, the forums were where the actual cult bonding happened. The parasocial relationship with Molyneux was the entry point, but the peer community on the forums was what made people hard to extract.

One super-participant wrote: "Certainty has been something I've craved since I was 10 and started reading Descartes' Meditations. Molyneux's theory finally got me there and it is as if someone has taken their boot off my chest and I can breathe for the first time." GNET That's Festinger in real time—the grand unified theory filling an emotional need, with the community then forming around the shared relief of certainty.

As cult expert Steven Hassan put it: "Molyneux knows how to talk like he knows what he's talking about, despite very little academic research." Southern Poverty Law Center

Sound familiar?

Never heard of it, but perfect! I now definitely smell an interesting article titled "Doomtubers!"

That's a strong title—punchy, immediately evocative, and it opens onto a genuinely rich set of ideas. You've already got the architecture of the piece sitting right here in this conversation.

The through-line is compelling: the ancient human appetite for apocalyptic narrative, how YouTube's incentive structure selects specifically for hedgehog epistemics, Festinger's insight that the prediction is really just social activation energy, the two-platform structure of cult formation, and the parasocial dynamic that both enables global scale and subtly dampens the deepest cohesion. That's a real argument with genuine movement through it, not just a collection of observations.

The Jiang case is a good entry point because he's more intellectually serious than the typical doomtuber—which actually makes the analysis sharper. It's easy to dismiss an obvious charlatan. It's more interesting to show how the medium's incentive structure corrupts even someone with real credentials and genuine historical knowledge. The platform doesn't just host the phenomenon, it breeds it.

And the Molyneux comparison gives you the darker end of the spectrum to triangulate against—where the same structural dynamics lead when there's no intellectual good faith underneath at all.

Festinger ties it together beautifully because he was essentially describing YouTube comment sections in 1956. That's your McLuhan move right there—the medium creates the conditions Festinger documented, at civilizational scale, continuously, algorithmically.

When it's ready, it'll fit right into the Banapana voice.

Jiang seems harmless at first—he's an academic, he's in front of a white board, he's well-spoken—and the supreme ingredient: he is really confident with no way to be disproven.

Yes—and that combination is almost a checklist for the phenomenon. The whiteboard is doing enormous work there. It's the visual grammar of legitimate pedagogy—it signals "this is a classroom, I am a teacher, what follows is knowledge." Combined with the academic affiliation, the Yale credential, the measured delivery, it creates what you might call an epistemic costume. The form is indistinguishable from genuine instruction.

The confidence without falsifiability is the key ingredient though, and it's worth being precise about why it's so effective. It's not just that he can't be disproven—it's that the framework is structured so that apparent disconfirmations get absorbed as further confirmation. The ceasefire happens and suddenly the prediction was always a "thought experiment about structural forces." That's not a bug in his presentation, it's the load-bearing feature. An audience that's already invested in the framework will do that reframing work themselves, enthusiastically, in the comments.

There's something almost predatory about the academic costume in this context—not necessarily consciously so, but functionally. The whiteboard and the measured tone lower exactly the defenses that might otherwise trigger skepticism. You don't interrogate a professor the way you'd interrogate a street corner prophet even if they're making equivalent epistemological claims.

And that's your McLuhan angle right there—the medium of the classroom, transplanted to YouTube, carries its authority signals intact while completely stripping away the accountability structures that made those signals meaningful in the first place. No peer review, no department, no colleagues pushing back, no obligation to update.

Just the whiteboard and a hundred thousand people nodding.

Awesome. Have at.

Let me read the Banapana editorial skill first before I write a word.

Viewed 2 files...

Viewed 2 files...

Good. Now I have everything I need. Let me write this.

Created a file, read a file...

Created a file, read a file...

There it is. About 1,400 words—tight enough to move, long enough to land the argument.

The structure follows the conversation's own logic: open on Jiang specifically, expose the epistemic costume, pivot to Tetlock's framework, then widen to the ancient appetite and YouTube's selection mechanism, bring in Festinger, use Molyneux as the darker end of the spectrum, and land on the medium/accountability point as the actual McLuhan thesis.

The ending earns its darkness—no false hope, just the whiteboard looking different with the sound off.

It's a first draft. Have at it with your red pen.

Claude's First Draft

Doomtubers

There's a man in Beijing with a whiteboard who knows exactly how civilization ends.

He'll tell you about it in a measured, academic tone—the kind that implies he has tenure somewhere important (he teaches at a progressive high school, but details). He'll draw arrows between historical events. He'll invoke game theory, Asimov, the Sicilian Expedition. He graduated from Yale. He's published. He has a Harvard affiliation. The whiteboard is spotless.

His name is Jiang Xueqin, he runs a YouTube channel called [Predictive History](#), and in May 2024 he recorded a classroom lecture predicting that Donald Trump would win the election and drag America into a catastrophic war with Iran. When those things started happening, the lecture went from a few hundred views to a hundred thousand subscribers in three days. The internet crowned him "China's Nostradamus." *Newsweek* wrote him up. He went viral in the most literal sense—something inert suddenly replicating everywhere.

The whiteboard did a lot of work in that lecture. It's the visual grammar of legitimate pedagogy—*this is a classroom, I am a teacher, what follows is knowledge*. Combined with the Yale credential, the measured delivery, the invocation of psychohistory (Isaac Asimov's fictional mathematical science of mass human behavior), and a framework that sounds rigorous without being falsifiable, you get something that looks indistinguishable from genuine insight. The epistemic costume is perfect.

Here's the thing: it's a costume.

The Model That Isn't

When Jiang refers to his charts as "models," something important collapses quietly and nobody notices. A chart showing historical patterns is a visualization. It can *illustrate* a hypothesis. But a model has formal structure—parameters, math, a

mechanism that generates specific predictions you can be *wrong about in advance*. Asimov's fictional psychohistory worked precisely because Hari Seldon could calculate probabilities on specific outcomes with quantified uncertainty. What Jiang does is the opposite: he uses the *aesthetic* of rigorous prediction—game theory vocabulary, historical pattern-matching, structural analysis—without the formal machinery that would make any of it testable.

Call it psychohistory cosplay.

The game theory framing has the same problem. Real game theory requires formally specifying players, available strategies, payoff functions, and information structures. When those are pinned down, you get specific equilibrium predictions you can be wrong about. But if "game theory" just means *I think about incentives narratively*, then it's not doing the predictive work—it's reasoned storytelling with prestige vocabulary stapled on.

Psychologist Philip Tetlock spent [decades running forecasting tournaments](#) with thousands of participants to figure out who actually predicts things well. His finding, borrowing Isaiah Berlin's old fox-and-hedgehog distinction: the people who know *one big thing* and organize their worldview around a grand unified theory—hedgehogs—perform barely better than chance. The people who draw on multiple frameworks, hold views with calibrated uncertainty, and update when evidence arrives—foxes—significantly outperform them.

Jiang is a hedgehog in a fox costume. The costume is very good.

The Broken Watch Problem

When the predicted events started materializing—Trump elected, U.S.-Iran tensions escalating—the clips went everywhere. What didn't go everywhere: the ceasefire Trump brokered, which was distinctly not the "boots on the ground, hostages not soldiers" catastrophe Jiang outlined. In the comments and forums,

this was smoothly reframed: the prediction was always a *thought experiment about structural forces*, not a literal forecast. The framework absorbed the contradiction without updating. This is not a bug. It is the [load-bearing feature](#).

One Reddit commenter, admirably resistant to the zeitgeist, noted that given the sheer number of self-proclaimed geopolitical forecasters online, one of them was statistically bound to get a scenario partially right. Even a broken watch is right twice a day. The internet finds that broken watch and makes it famous, and then the watch is no longer "broken"—it's *prescient*. The watches that were wrong at the same moment are forgotten entirely.

This is not how we should be evaluating forecasters, and Tetlock would tell you so at length.

[The Free Press](#) also noted, almost as a footnote, that Jiang believes a coalition of Freemasons, Jesuits, and members of a defunct Jewish cult are conspiring to rule the world from Jerusalem—which is either a red flag about the framework underneath, or proof that the whiteboard really is doing enormous work.)

Ancient Appetite, New Distribution

Here's what makes this more than a story about one YouTube channel: the appetite Jiang is feeding is ancient and remarkably stable across cultures. Every civilization generates its own eschatological framework. The Norse had Ragnarök. The Zoroastrians had Frashokereti. Christians have Revelation. We have geopolitical YouTube. The packaging changes. The demand doesn't.

Apocalyptic narratives are [paradoxically comforting](#)—they impose meaning and structure on what otherwise feels like chaotic, uncontrollable events. If civilization is collapsing according to a *plan*, at least someone understands it. The alternative—that history is messy contingency with no grand arc and no hidden architect—is

genuinely harder to sit with. Certainty, even catastrophic certainty, is preferable to the fox's maddening "it depends."

There's also a status dimension. Knowing about the coming catastrophe before everyone else feels like insider knowledge. You're not a passive victim of events—you're one of the few who see. This is the same psychological reward structure as any conspiracy theory, dressed in a blazer and standing in front of a whiteboard.

YouTube's incentive structure selects hard for exactly these qualities. Confident, vivid, high-stakes narratives get clicks. Calibrated uncertainty gets skipped. A thumbnail reading "WHY THE EMPIRE IS ABOUT TO COLLAPSE" outperforms "here are some probabilistic tendencies worth monitoring" by an order of magnitude. The platform doesn't just host doomtubers—it *breeds* them, by systematically rewarding the epistemic vices Tetlock identified and penalizing the epistemic virtues. The recommendation algorithm then clusters them together, so viewers who watch one get fed the entire ecosystem, progressively less tolerant of ambiguity.

McLuhan would have recognized this immediately: the medium doesn't just carry the content, it *selects* for it. YouTube is a hedgehog factory.

Festinger's Prophecy

In 1956, psychologist Leon Festinger [infiltrated a UFO doomsday cult](#) in Chicago that had predicted a great flood would destroy the world on a specific date. When the date passed uneventfully, the group didn't dissolve—it intensified. They began proselytizing harder than ever. Festinger developed his theory of cognitive dissonance to explain it: the greater the investment in a belief, the more creative the mind becomes in protecting it from disconfirming evidence. The prediction failing wasn't experienced as refutation—it was reframed as confirmation that the group's faith had *saved* the world.

But here's the part worth sitting with: the prediction was almost a pretext. It provided the initial activation energy for people to find each other and form bonds. Once those bonds existed, the social reality of the group *became* the thing being defended—not the original claim. The theology was fungible. The belonging wasn't.

Festinger essentially described YouTube comment sections in 1956.

The doomtuber phenomenon has a structural wrinkle Festinger couldn't have anticipated, though: the parasocial nature of the YouTube relationship partially dampens the social cohesion. You have a one-sided relationship with the lecturer—he doesn't know you exist—which limits the depth of community that can form around the channel itself. What actually happens, in the documented cases, is a two-platform structure: YouTube as recruitment funnel, with Discord servers, Subreddits, and private forums as the layer where genuine bonding occurs. The parasocial relationship is the entry point; the peer community is what makes extraction difficult.

Stefan Molyneux's [Freedomain Radio](#) is the clearest example of where this architecture leads when the intellectual good faith runs out entirely. YouTube channel, private forums, a framework of pseudo-therapeutic certainty, and a practice he invented called "deFOOing"—departing from one's Family of Origin—that functioned as the community's loyalty test. Multiple cult experts identified the dynamics. The British Cult Information Centre called it a cult. When the deFOOing started making international news, Molyneux's response was: "If I advised a wife to leave an abusive husband, there would not be articles about how I am a cult leader." The whiteboard equivalent of the thought experiment defense.

The Costume Is the Message

What makes Jiang more interesting than Molyneux—and more instructive—is that he's genuinely more intellectually serious. The Yale education, the engagement

with actual history and philosophy, the classroom context—those are real. He has read the books. He knows the arguments. He's not purely performing expertise; he has some.

Which is exactly what makes the epistemic costume so effective, and so worth examining. It's easy to dismiss an obvious charlatan. It's harder—and more important—to articulate why a credentialed, well-read, thoughtful person, operating in the medium of the classroom, produces the same audience dynamic as a prophet with a sandwich board. The answer isn't that Jiang is secretly a charlatan. The answer is that the platform transforms what the classroom *does*.

In an actual institution, the whiteboard comes with accountability structures: peer review, departmental colleagues, students who challenge you, an obligation to update when evidence contradicts your framework. The classroom's epistemic authority signals are meaningful because they're backed by those structures. Transplanted to YouTube, the authority signals survive intact. The accountability structures do not. You keep the whiteboard. You lose everything that made the whiteboard trustworthy.

The medium strips the accountability while preserving the aesthetic. A hundred thousand people watch a professor who answers to no one, in a classroom that never ends, teaching a subject that cannot be graded.

Festinger would recognize the room. He just wouldn't have predicted the scale.

If you want to understand how the platform shapes the prophet, [Tetlock's Superforecasting](#) is where to start—and then come back and watch any doomtuber with the sound off. The whiteboard will look different.